WHITE / GREEN PAPER SERIES

AI Disclosure Events (AIDE): Internal Briefing

Schema v0.1 + Pilot Plan

Executive Summary

This briefing proposes a minimal, machine■actionable AI Disclosure Event (AIDE) schema to make AI use in research transparent, citable, and auditable. AIDE leverages existing PID rails—ORCID, DataCite DOIs, ROR, and RAiD—to capture who/when/what/why of substantive AI use, while supporting privacy via hashed prompt/output capsules.

Schema Overview (v0.1)

Core fields: id (DOI), created (timestamp), actor (ORCID, ROR affiliation), context (RAiD, grant, related outputs), usage (category, purpose, role, stage), model (name, version, provider, model PID, hostingOrg), data (dataset/source DOIs), prompts (promptClass, hashes), oversight (humanInTheLoop, responsibleUse, notes). See attached JSON Schema and XSD files for validation details.

Implementation & Resolution

AIDE records can be minted as DataCite DOIs with a human■readable landing page and JSON API, and/or pushed to ORCID as a contribution. Resolvers can expose a compact Provenance Badge for PDFs/preprints showing key fields and linking to the full disclosure.

Pilots (3–6 months)

1) Publisher submission capture and display; 2) Repository capture with ORCID push; 3) Model release DOIs referencing training data DOIs. Deliverables: schema v0.1, two pilot reports, badge spec v0.1.

Artifacts

• aide_schema_v0_1.json (JSON Schema) • aide_schema_v0_1.xsd (XML Schema) • ddi_aide_profile_notes.txt (DDI profile mapping) • aide_example_chatgpt.json / .xml • aide_example_claude.json / .xml • aide_example_huggingface.json / .xml